

HITESH BISHT

MLOps Engineer

@ hbsht1110@gmail.com +91-7668064366 in hitesh-bisht hiteshbisht11

SUMMARY

MLOps Engineer with 2+ years of experience building and operating scalable ML systems for high-throughput inference. Skilled in model deployment, monitoring, and optimization, with a focus on reliability, observability, and efficient production ML pipelines.

EXPERIENCE

MLOps Engineer Bangalore, India
Zzazz.ai

Aug 2025 – Present

Observability Platform:

- Built an end-to-end observability platform for 30+ ML services using Prometheus, Grafana, and OpenTelemetry.
- Instrumented request count, failure rate, API & upstream latency; added structured logs and Grafana dashboards, reducing MTTR for production incidents.
- Added per-step latency tracing via OpenTelemetry to surface bottlenecks; set Alertmanager thresholds on critical inference paths.
- Built a shadow deployment framework to compare candidate vs. production model outputs, alerting on significant deviations for zero-downtime rollouts.

Model Microservices Platform:

- Led end-to-end deployment of 30+ ML models serving ~500K inference requests per day across multiple internal teams.
- Containerized models using Docker & Kubernetes, achieving 99.5% service availability and enabling horizontal auto-scaling from 2 → 20 pods based on traffic spikes.
- Reduced inference latency by 75% (800ms → 200ms avg response time) through optimized CPU/memory allocation and HPA tuning.

Signal Automation:

- Message-driven pipeline (FastAPI + RabbitMQ) processing 50K+ publisher URLs daily.
- Built an ML pipeline (scikit-learn, DistilUSE ONNX, TensorFlow) that auto-generates DOM-path CSS selector rules for pricing widget placement on publisher websites.
- Crawled publisher homepages via Selenium + NetworkX to extract container paths, cutting integration time from 3–5 days to under 6 hours (85% reduction).
- Migrated inference workloads from GPU to optimized CPU instances, reducing infrastructure cost by 42% annually (~INR 18–22 Lakhs saved per year) while maintaining SLA.

Software Engineer Noida, India
Gauge Data Solutions Pvt Ltd

Oct 2024 – Aug 2025

Crawlers:

- Scaled data pipelines aggregating from 20+ court portals, processing 50K–70K documents weekly (~3M annually).
- Built a dedicated crawler for Supreme Court High Court portals, extracting judgements and cause lists for downstream legal analytics.
- Improved scraping success rate from 78% to 94% by implementing retry/backoff strategies.

Microservices & APIs:

- Developed REST APIs supporting 100K+ internal requests per day.
- Reduced data validation errors by 35% through schema enforcement and automated checks.

Software Engineer (Intern) Noida (Remote)
KaveoSoft Innovations Pvt Ltd

Nov 2023 – Sept 2024

EFT Transaction:

- Developed and maintained EFT transaction processing modules handling 10K+ daily transactions.
- Designed and deployed a Spring Boot REST API for EFT transaction.
- Enhanced transaction validation accuracy by 25%, reducing reconciliation issues.
- Maintained zero high-severity production incidents during tenure.

EDUCATION

B.Tech in CSE : 8.5 CGPA
Graphic Era Hill University

2020 – 2024 Bhimtal, Uttarakhand

TECHNICAL SKILLS

- Programming Languages**
Python, Java, PL/SQL
- Databases**
PostgreSQL, MongoDB, Redis, Elasticsearch, Oracle
- ML & Data**
PyTorch, TensorFlow, Scikit-learn, Pandas, NumPy
- Containerization/orchestration**
Docker, Kubernetes*
- CI/CD**
GitLab Pipelines, Jenkins*
- Workflow Orchestration**
Airflow

*Elementary Proficiency

PROJECTS

Real-Time Financial Risk Intelligence Platform ([GitHub](#)):

- Real-time transaction risk scoring API processing 1K–2K transactions/min with p95 latency under 150ms.
- Fraud detection model (LightGBM/XGBoost) on 750K+ synthetic & historical transactions (2.1% fraud rate), achieving PR-AUC: 0.89, Recall: 78% at 5% false positive rate.
- Real-time behavioral features (transaction velocity, spend deviation, merchant risk score) with feature computation latency < 25ms per request.
- Online feature store using Redis + Feast, supporting 200+ RPS sustained inference throughput.
- Graph-based risk propagation using NetworkX, modeling 50K+ entity relationships, improving fraud recall by 10%.
- Tech Stack:** Python, LightGBM, scikit-learn, Kafka (Redpanda), Redis, Feast (Feature Store), FastAPI, Docker, MLflow, Airflow, NetworkX, Prometheus & Grafana

CERTIFICATIONS:

- Database Management Essentials ([link](#))